(/)

You have a captain's log due before 2020-06-07 (in 1 day)! Log it now! (/captain_logs/35487/edit)

0x00. Binary Classification

- Specializations Machine Learning Supervised Learning
- ▲ by Alexa Orrico, Software Engineer at Holberton School
- 🗱 weight: 3
- Project over took place from 02-10-2020 to 02-19-2020 you're done with 79% of tasks.
- ☑ Manual QA review was done by Abdel Giovanny Perez on 02-25-2020.
- QA review fully automated.

In a nutshell...

- Manual QA review: 8.5/42 mandatory
 Auto QA review: 142.0/161 mandatory
- Altogether: 74.14%
 - o Mandatory: 74.14%
 - o Optional: no optional tasks





Background Context

Welcome to your first project on supervised learning! At the end of this project, you should be able to build your own binary image classifier from scratch using <code>numpy</code>. As you might already see, there are a <code>LOT</code> of resources for you to read/watch. It may be tempting to dive into the projects right away, but it is <code>HIGHLY RECOMMENDED</code> that you spend <code>AT LEAST 1</code> whole day going over the following materials. You should only start the project once you have a decent understanding of all the topics mentioned in <code>Learning Objectives</code>. You may also notice that there are multiple resources that cover the same topic, with some more technical than others. If you find yourself getting lost in a resource, move on to another and come back to the more technical one after you intuitively understand that topic. Good luck and have fun!

Resources

Read or watch:

- Supervised vs. Unsupervised Machine Learning (/rltoken/16mL gwlZqRa3NAv325YiQ)
- How would you explain neural networks to someone who knows very little about AI or neurology? (/rltoken/A1B1qLJgpKvA4RySpI6r7g)
- Using Neural Nets to Recognize Handwritten Digits (/rltoken/EjjENEVXJKiAZsqSqWMI-w) (until "A simple network to classify handwritten digits" (excluded))
- Forward propagation (/rltoken/da0lo6JbjEDxCbffjRHc7g)
- Understanding Activation Functions in Neural Networks (/rltoken/hz77ChKoiSjMFzi7mgMzBA)
- Loss function (/rltoken/KgRV0-l2LBdQciUXGleCqQ)
- Gradient descent (/rltoken/7iSJelYELwy7C8cCsGk5hw)
- Calculus on Computational Graphs: Backpropagation (/rltoken/BONZS65eZnlMjngFhr7dPA)
- Backpropagation calculus (/rltoken/Arpa6EFk9q5gD9aJ4Wl5qA)
- What is a Neural Network? (/rltoken/EkncpxTwCUJztJsYI0wciA)
- Supervised Learning with a Neural Network (/rltoken/LoWxJZN-JA0VkV13QQrw1g)
- Binary Classification (/rltoken/cFuQ0hUHg_SpVCHvKrXMzg)
- Logistic Regression (/rltoken/sGllY030fFNX4nNQidq5fw)
- Logistic Regression Cost Function (/rltoken/xZTVYTU5pSnSK3a7o3OulQ)
- Gradient Descent (/rltoken/M3YbEr_BqYcILJNz7YzLaQ)
- Computation Graph (/rltoken/X5CelY1ajZt3wHrjtls6Fg)
- Logistic Regression Gradient Descent (/rltoken/HLxYo6tgVumVNRysPUxKNA)
- Vectorization (/rltoken/Zxdbe_-GWZRfXKfs5JM9ig)
- Vectorizing Logistic Regression (/rltoken/HQ9VuO9c4XPJglm6Nxyn5Q)
- Vectorizing Logistic Regression's Gradient Computation (/rltoken/RaswXJ2G9LHswV0CypjA0A)
- A Note on Python/Numpy Vectors (/rltoken/wKRb7J-yeA92EF5aEJd3oA)
- Neural Network Representations (/rltoken/JyhRr98YlhACYERu0GncNQ)
- Computing Neural Network Output (/rltoken/Lpcj3uH 6hh8Fp1dXzKkCw)
- Vectorizing Across Multiple Examples (/rltoken/uWY4JFKkT58mrHSBig f5A)
- Gradient Descent For Neural Networks (/rltoken/Q583jTfE8BfU5hPW1xlDiQ)
- Random Initialization (/rltoken/2uZWU7WaWSQfwGNKlcgZig)
- Deep L-Layer Neural Network (/rltoken/iMyExMqGZGcawyK51FcdzQ)
- Train/Dev/Test Sets (/rltoken/varxWT03Dy39WlyrZBGAVQ)
- Random Initialization For Neural Networks: A Thing Of The Past (/rltoken/gMbQAlqJulDEppMx5rESBQ)
- Initialization of deep networks (/rltoken/ymVrn0IFwxnzoo3WwEZzcQ)
- numpy.zeros (/rltoken/Ho3q2XEingriAApbvjTy6w)
- numpy.random.randn (/rltoken/JFRI j ebMhXasgQ geUug)
- numpy.exp (/rltoken/AjRCO7Yh0JW4lC0xuGmbOg)
- numpy.log (/rltoken/AtD0YS-8Q3Oc0TN1MQ6Wuw)
- numpy.sqrt (/rltoken/Xb26cP7K6Bqtoc NhQXosg)
- numpy.where (/rltoken/6owcKbc6MWk7JlnlM0AaWA)

Optional:

- Predictive analytics (/rltoken/Bp 21TFndH5cSWLBObZuzQ)
- Maximum Likelihood Estimation (/rltoken/1gmRZd9vF7l3bclPMnTPaw)

Learning Objectives

At the end of this project, you are expected to be able to explain to anyone (/rltoken/yKtle6fSEklenpymZ18Xrw), without the help of Google:

General

- What is a model?
- · What is supervised learning?
- What is a prediction?
- What is a node?
- What is a weight?
- · What is a bias?
- What are activation functions?
 - o Sigmoid?
 - o Tanh?
 - o Relu?
 - Softmax?
- · What is a layer?
- What is a hidden layer?
- What is Logistic Regression?
- What is a loss function?
- What is a cost function?
- What is forward propagation?
- What is Gradient Descent?
- What is back propagation?
- What is a Computation Graph?
- How to initialize weights/biases
- The importance of vectorization
- · How to split up your data

Requirements

General

- Allowed editors: vi, vim, emacs
- All your files will be interpreted/compiled on Ubuntu 16.04 LTS using python3 (version 3.5)
- Your files will be executed with numpy (version 1.15)
- · All your files should end with a new line
- The first line of all your files should be exactly #!/usr/bin/env python3
- A README.md file, at the root of the folder of the project, is mandatory
- Your code should use the pycodestyle style (version 2.4)
- All your modules should have documentation (python3 -c

```
'print(__import__("my_module").__doc__)')
```

- All your classes should have documentation (python3 -c
 'print(__import__("my_module").MyClass.__doc__)')
- All your functions (inside and outside a class) should have documentation (python3 -c 'print(__import__("my_module").my_function.__doc__)' and python3 -c 'print(__import__("my_module").MyClass.my_function.__doc__)')
- Unless otherwise noted, you are not allowed to import any module except import numpy as np
- Unless otherwise noted, you are not allowed to use any loops (for, while, etc.)
- All your files must be executable
- The length of your files will be tested using wc

More Info

Matrix Multiplications

For all matrix multiplications in the following tasks, please use numpy.matmul (/rltoken/Ox8bY8ogmUftzjR96lrMDw)

Testing your code

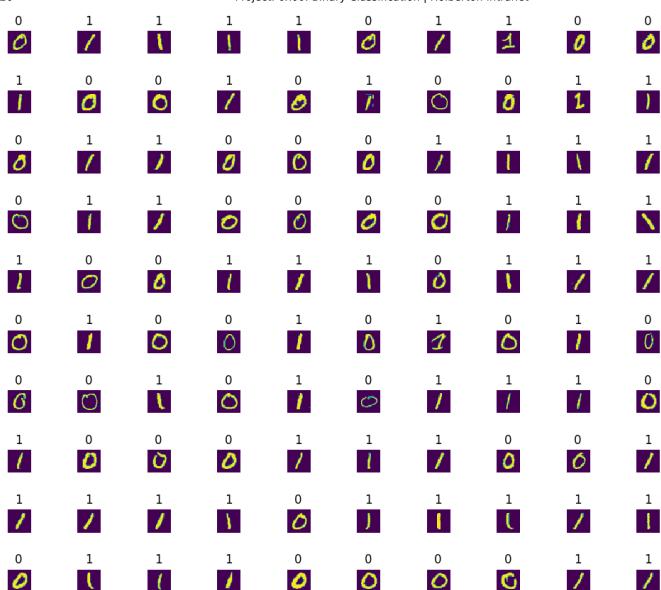
In order to test your code, you'll need DATA! Please download these datasets (Binary_Train.npz (https://s3.amazonaws.com/intranet-projects-files/holbertonschool-ml/Binary_Train.npz), Binary_Dev.npz (https://s3.amazonaws.com/intranet-projects-files/holbertonschool-ml/Binary_Dev.npz)) to go along with all of the following main files. You **do not** need to upload these files to GitHub. Your code will not necessarily be tested with these datasets. All of the following code assumes that you have stored all of your datasets in a separate data directory.

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat show_data.py
#!/usr/bin/env python3

import matplotlib.pyplot as plt
import numpy as np

lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']

fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_3D[i])
    plt.title(Y[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./show_data.py
```



Tasks

O. Neuron mandatory

Score: 100.00% (Checks completed: 100.00%)

Write a class Neuron that defines a single neuron performing binary classification:

• class constructor: def __init__(self, nx):



- o nx is the number of input features to the neuron
 - If nx is not an integer, raise a TypeError with the exception: nx must be an integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be a positive integer
- All exceptions should be raised in the order listed above
- · Public instance attributes:
 - W: The weights vector for the neuron. Upon instantiation, it should be initialized using a random normal distribution.
 - o b: The bias for the neuron. Upon instantiation, it should be initialized to 0.
 - A: The activated output of the neuron (prediction). Upon instantiation, it should be initialized to 0.

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 0-main.py
#!/usr/bin/env python3
import numpy as np
Neuron = __import__('0-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X.shape[0])
print(neuron.W)
print(neuron.W.shape)
print(neuron.b)
print(neuron.A)
neuron.A = 10
print(neuron.A)
alexa@ubuntu-xenial:0x00-binary_classification$ ./0-main.py
1.86755799e+00 -9.77277880e-01 9.50088418e-01 -1.51357208e-01
. . .
  -5.85865511e-02 -3.17543094e-01 -1.63242330e+00 -6.71341546e-02
  1.48935596e+00 5.21303748e-01 6.11927193e-01 -1.34149673e+00
(1, 784)
0
0
10
alexa@ubuntu-xenial:0x00-binary_classification$
```

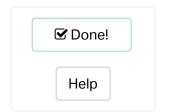
- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 0-neuron.py

Check your code?

QA Review

1. Privatize Neuron mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class Neuron that defines a single neuron performing binary classification (Based on 0-neuron.py):

- class constructor: def __init__(self, nx):
 - o nx is the number of input features to the neuron
 - If nx is not an integer, raise a TypeError with the exception: nx must be a integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be positive
 - All exceptions should be raised in the order listed above
- Private instance attributes:
 - __w: The weights vector for the neuron. Upon instantiation, it should be initialized using a random normal distribution.
 - __b : The bias for the neuron. Upon instantiation, it should be initialized to 0.
 - __A: The activated output of the neuron (prediction). Upon instantiation, it should be initialized to
 - Each private attribute should have a corresponding getter function (no setter function).

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 1-main.py
#!/usr/bin/env python3
import numpy as np
Neuron = __import__('1-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X.shape[0])
print(neuron.W)
print(neuron.b)
print(neuron.A)
neuron.A = 10
print(neuron.A)
alexa@ubuntu-xenial:0x00-binary_classification$ ./1-main.py
1.86755799e+00 -9.77277880e-01 9.50088418e-01 -1.51357208e-01
. . .
  -5.85865511e-02 -3.17543094e-01 -1.63242330e+00 -6.71341546e-02
  1.48935596e+00 5.21303748e-01 6.11927193e-01 -1.34149673e+00
0
0
Traceback (most recent call last):
 File "./1-main.py", line 16, in <module>
   neuron.A = 10
AttributeError: can't set attribute
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 1-neuron.py

Check your code? QA Review

2. Neuron Forward Propagation

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class Neuron that defines a single neuron performing binary classification (Based on 1-neuron.py):

- Add the public method def forward_prop(self, X):
 - Calculates the forward propagation of the neuron
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Updates the private attribute ___A
 - The neuron should use a sigmoid activation function
 - Returns the private attribute __A

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 2-main.py
#!/usr/bin/env python3
import numpy as np
Neuron = __import__('2-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X.shape[0])
neuron._Neuron__b = 1
A = neuron.forward_prop(X)
if (A is neuron.A):
        print(A)
vagrant@ubuntu-xe
alexa@ubuntu-xenial:0x00-binary_classification$ ./2-main.py
[[5.34775247e-10 7.24627778e-04 4.52416436e-07 ... 8.75691930e-05
  1.13141966e-06 6.55799932e-01]]
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 2-neuron.py

QA Review Check your code?

3. Neuron Cost | mandatory

Score: 100.00% (Checks completed: 100.00%)

Write a class Neuron that defines a single neuron performing binary classification (Based on 2-neuron.py):

☑ Done!

- Add the public method def cost(self, Y, A):
 - Calculates the cost of the model using logistic regression
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - A is a numpy.ndarray with shape (1, m) containing the activated output of the neuron for each example
 - o To avoid division by zero errors, please use 1.0000001 A instead of 1 A
 - o Returns the cost

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 3-main.py
#!/usr/bin/env python3
import numpy as np

Neuron = __import__('3-neuron').Neuron

lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T

np.random.seed(0)
neuron = Neuron(X.shape[0])
A = neuron.forward_prop(X)
cost = neuron.cost(Y, A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./3-main.py
4.365104944262272
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 3-neuron.py

Check your code? QA Review

4. Evaluate Neuron

mandatory

Score: 100.00% (Checks completed: 100.00%)

Write a class Neuron that defines a single neuron performing binary classification (Based on 3-neuron.py):



- Add the public method def evaluate(self, X, Y):
 - Evaluates the neuron's predictions
 - X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - Returns the neuron's prediction and the cost of the network, respectively
 - The prediction should be a numpy.ndarray with shape (1, m) containing the predicted labels for each example
 - The label values should be 1 if the output of the network is >= 0.5 and 0 otherwise

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 4-main.py
#!/usr/bin/env python3
import numpy as np
Neuron = __import__('4-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X.shape[0])
A, cost = neuron.evaluate(X, Y)
print(A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./4-main.py
[[0 0 0 ... 0 0 0]]
4.365104944262272
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 4-neuron.py

QA Review Check your code?

5. Neuron Gradient Descent | mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class Neuron that defines a single neuron performing binary classification (Based on 4-neuron.py):

- Add the public method def gradient_descent(self, X, Y, A, alpha=0.05):
 - Calculates one pass of gradient descent on the neuron
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - A is a numpy.ndarray with shape (1, m) containing the activated output of the neuron for each example
 - o alpha is the learning rate
 - Updates the private attributes __w and __b

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 5-main.py
#!/usr/bin/env python3
import numpy as np
Neuron = __import__('5-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X.shape[0])
A = neuron.forward_prop(X)
neuron.gradient_descent(X, Y, A, 0.5)
print(neuron.W)
print(neuron.b)
alexa@ubuntu-xenial:0x00-binary_classification$ ./5-main.py
1.86755799e+00 -9.77277880e-01 9.50088418e-01 -1.51357208e-01
. . .
  -5.85865511e-02 -3.17543094e-01 -1.63242330e+00 -6.71341546e-02
  1.48935596e+00 5.21303748e-01 6.11927193e-01 -1.34149673e+00]
0.2579495783615682
alexa@ubuntu-xenial:0x00-binary_classification$
```

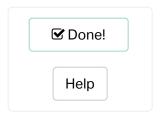
Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 5-neuron.py

Check your code? QA Review

6. Train Neuron mandatory

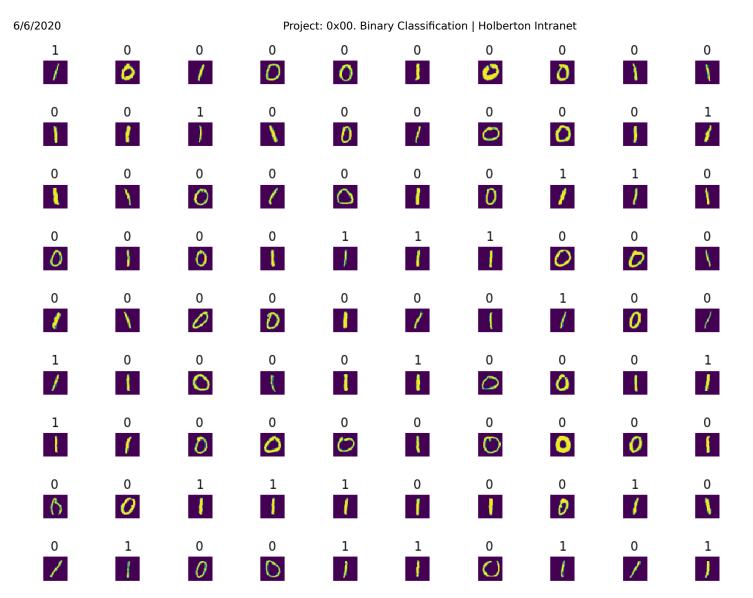
Score: 100.00% (Checks completed: 100.00%)



Write a class Neuron that defines a single neuron performing binary classification (Based on 5-neuron.py):

- Add the public method def train(self, X, Y, iterations=5000, alpha=0.05):
 - Trains the neuron
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
 - alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
 - All exceptions should be raised in the order listed above
 - Updates the private attributes __w , __b , and __A
 - You are allowed to use one loop
 - Returns the evaluation of the training data after iterations of training have occurred

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 6-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
Neuron = __import__('6-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X_train.shape[0])
A, cost = neuron.train(X_train, Y_train, iterations=10)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = neuron.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./6-main.py
Train cost: 1.3805076999077135
Train accuracy: 64.73746545598105%
Dev cost: 1.4096194345468178
Dev accuracy: 64.49172576832152%
```



Not that great... Let's get more data!

1

0

Repo:

0

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification

1

0

O

1

0

O

• File: 6-neuron.py

Check your code? QA Review

7. Upgrade Train Neuron

mandatory

Score: 0.00% (Checks completed: 0.00%)

Q

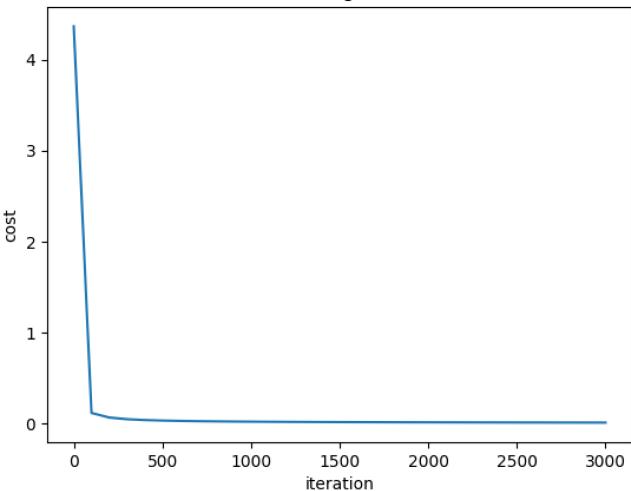
0

Write a class Neuron that defines a single neuron performing binary classification (Based on 6-neuron.py):

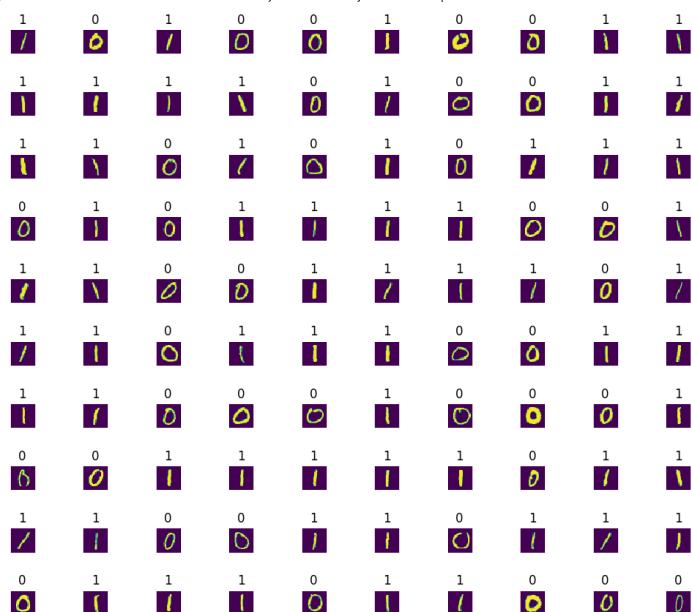
- Update the public method train to def train(self, X, Y, iterations=5000, alpha=0.05, verbose=True, graph=True, step=100):
- □ Done?
- Trains the neuron by updating the private attributes __w , __b , and __A
- X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
- Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
- o iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
- alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
- verbose is a boolean that defines whether or not to print information about the training. If True,
 print Cost after {iteration} iterations: {cost} every step iterations:
 - Include data from the 0th and last iteration
- graph is a boolean that defines whether or not to graph information about the training once the training has completed. If True:
 - Plot the training data every step iterations as a blue line
 - Label the x-axis as iteration
 - Label the y-axis as cost
 - Title the plot Training Cost
 - Include data from the 0th and last iteration
- o Only if either verbose or graph are True:
 - if step is not an integer, raise a TypeError with the exception step must be an integer
 - if step is not positive or is greater than iterations, raise a ValueError with the exception step must be positive and <= iterations</p>
- All exceptions should be raised in the order listed above
- The 0th iteration should represent the state of the neuron before any training has occurred
- You are allowed to use one loop
- You can use import matplotlib.pyplot as plt
- Returns: the evaluation of the training data after iterations of training have occurred

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 7-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
Neuron = __import__('7-neuron').Neuron
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
neuron = Neuron(X_train.shape[0])
A, cost = neuron.train(X_train, Y_train, iterations=3000)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = neuron.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./7-main.py
Cost after 0 iterations: 4.365104944262272
Cost after 100 iterations: 0.11955134491351888
Cost after 3000 iterations: 0.013386353289868338
```





Train cost: 0.013386353289868338 Train accuracy: 99.66837741808132% Dev cost: 0.010803484515167197 Dev accuracy: 99.81087470449172%



- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 7-neuron.py



8. NeuralNetwork mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification:

- class constructor: def __init__(self, nx, nodes):
 - o nx is the number of input features
 - If nx is not an integer, raise a TypeError with the exception: nx must be an integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be a positive integer
 - nodes is the number of nodes found in the hidden layer
 - If nodes is not an integer, raise a TypeError with the exception: nodes must be an integer
 - If nodes is less than 1, raise a ValueError with the exception: nodes must be a positive integer
 - All exceptions should be raised in the order listed above
- Public instance attributes:
 - W1: The weights vector for the hidden layer. Upon instantiation, it should be initialized using a random normal distribution.
 - o b1: The bias for the hidden layer. Upon instantiation, it should be initialized with 0's.
 - o A1: The activated output for the hidden layer. Upon instantiation, it should be initialized to 0.
 - W2: The weights vector for the output neuron. Upon instantiation, it should be initialized using a random normal distribution.
 - o b2: The bias for the output neuron. Upon instantiation, it should be initialized to 0.
 - A2: The activated output for the output neuron (prediction). Upon instantiation, it should be initialized to 0.

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 8-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('8-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
print(nn.W1)
print(nn.W1.shape)
print(nn.b1)
print(nn.W2)
print(nn.W2.shape)
print(nn.b2)
print(nn.A1)
print(nn.A2)
nn.A1 = 10
print(nn.A1)
alexa@ubuntu-xenial:0x00-binary_classification$ ./8-main.py
[[ 1.76405235  0.40015721  0.97873798  ...  0.52130375  0.61192719
  -1.34149673]
 \begin{bmatrix} 0.47689837 & 0.14844958 & 0.52904524 & \dots & 0.0960042 & -0.0451133 \end{bmatrix}
   0.07912172]
 [ 0.85053068 -0.83912419 -1.01177408 ... -0.07223876  0.31112445
  -1.07836109]]
(3, 784)
[[0.]
 [0.]
 [0.]]
[[ 1.06160017 -1.18488744 -1.80525169]]
(1, 3)
0
0
0
10
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 8-neural_network.py

Check your code? QA Review

9. Privatize NeuralNetwork

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 8-neural_network.py):

- class constructor: def __init__(self, nx, nodes):
 - nx is the number of input features
 - If nx is not an integer, raise a TypeError with the exception: nx must be an integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be a positive integer
 - nodes is the number of nodes found in the hidden layer
 - If nodes is not an integer, raise a TypeError with the exception: nodes must be an integer
 - If nodes is less than 1, raise a ValueError with the exception: nodes must be a positive integer
 - o All exceptions should be raised in the order listed above
- Private instance attributes:
 - W1: The weights vector for the hidden layer. Upon instantiation, it should be initialized using a random normal distribution.
 - o b1: The bias for the hidden layer. Upon instantiation, it should be initialized with 0's.
 - A1: The activated output for the hidden layer. Upon instantiation, it should be initialized to 0.
 - W2: The weights vector for the output neuron. Upon instantiation, it should be initialized using a random normal distribution.
 - o b2: The bias for the output neuron. Upon instantiation, it should be initialized to 0.
 - A2: The activated output for the output neuron (prediction). Upon instantiation, it should be initialized to 0.
 - Each private attribute should have a corresponding getter function (no setter function).

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 9-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('9-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
print(nn.W1)
print(nn.b1)
print(nn.W2)
print(nn.b2)
print(nn.A1)
print(nn.A2)
nn.A1 = 10
print(nn.A1)
alexa@ubuntu-xenial:0x00-binary_classification$ ./9-main.py
[[ 1.76405235  0.40015721  0.97873798  ...  0.52130375  0.61192719
  -1.34149673]
 [ 0.47689837  0.14844958  0.52904524  ...  0.0960042
                                                      -0.0451133
   0.07912172]
 [ 0.85053068 -0.83912419 -1.01177408 ... -0.07223876  0.31112445
  -1.07836109]]
[[0.]
 [0.]
 [0.]]
[[ 1.06160017 -1.18488744 -1.80525169]]
0
Traceback (most recent call last):
  File "./9-main.py", line 19, in <module>
    nn.A1 = 10
AttributeError: can't set attribute
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 9-neural_network.py

Check your code?

QA Review

10. NeuralNetwork Forward Propagation

Score: 100.00% (Checks completed: 100.00%)

mandatory

☑ Done!

Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 9-neural_network.py):

- Add the public method def forward_prop(self, X):
 - Calculates the forward propagation of the neural network
 - X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Updates the private attributes __A1 and __A2
 - The neurons should use a sigmoid activation function
 - Returns the private attributes __A1 and __A2 , respectively

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 10-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('10-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
nn._NeuralNetwork__b1 = np.ones((3, 1))
nn._NeuralNetwork__b2 = 1
A1, A2 = nn.forward_prop(X)
if A1 is nn.A1:
        print(A1)
if A2 is nn.A2:
        print(A2)
alexa@ubuntu-xenial:0x00-binary_classification$ ./10-main.py
[[5.34775247e-10 7.24627778e-04 4.52416436e-07 ... 8.75691930e-05
  1.13141966e-06 6.55799932e-01]
 [9.99652394e-01 9.99999995e-01 6.77919152e-01 ... 1.00000000e+00
  9.99662771e-01 9.99990554e-01]
 [5.57969669e-01 2.51645047e-02 4.04250047e-04 ... 1.57024117e-01
  9.97325173e-01 7.41310459e-02]]
 [[0.23294587 \ 0.44286405 \ 0.54884691 \ \dots \ 0.38502756 \ 0.12079644 \ 0.593269 ] 
                                                                         11
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

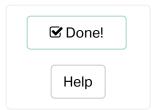
- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification

File: 10-neural_network.py

QA Review Check your code?

11. NeuralNetwork Cost mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 10-neural_network.py):

- Add the public method def cost(self, Y, A):
 - Calculates the cost of the model using logistic regression
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - A is a numpy.ndarray with shape (1, m) containing the activated output of the neuron for each example
 - To avoid division by zero errors, please use 1.0000001 A instead of 1 A
 - Returns the cost

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 11-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('11-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
_, A = nn.forward_prop(X)
cost = nn.cost(Y, A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./11-main.py
0.7917984405648548
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 11-neural_network.py

Check your code?

QA Review

12. Evaluate NeuralNetwork

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 11-neural_network.py):

- Add the public method def evaluate(self, X, Y):
 - Evaluates the neural network's predictions
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - Returns the neuron's prediction and the cost of the network, respectively
 - The prediction should be a numpy.ndarray with shape (1, m) containing the predicted labels for each example
 - The label values should be 1 if the output of the network is >= 0.5 and 0 otherwise

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 12-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('12-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
A, cost = nn.evaluate(X, Y)
print(A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./12-main.py
[[0 0 0 ... 0 0 0]]
0.7917984405648548
alexa@ubuntu-xenial:0x00-binary_classification$
```

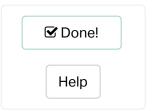
Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 12-neural_network.py

Check your code? **QA Review**

13. NeuralNetwork Gradient Descent mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 12-neural_network.py):

- Add the public method def gradient_descent(self, X, Y, A1, A2, alpha=0.05):
 - Calculates one pass of gradient descent on the neural network
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - A1 is the output of the hidden layer
 - A2 is the predicted output
 - alpha is the learning rate
 - Updates the private attributes __w1 , __b1 , __w2 , and __b2

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 13-main.py
#!/usr/bin/env python3
import numpy as np
NN = __import__('13-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
nn = NN(X.shape[0], 3)
A1, A2 = nn.forward_prop(X)
nn.gradient_descent(X, Y, A1, A2, 0.5)
print(nn.W1)
print(nn.b1)
print(nn.W2)
print(nn.b2)
alexa@ubuntu-xenial:0x00-binary_classification$ ./13-main.py
[[ 1.76405235  0.40015721  0.97873798  ...  0.52130375  0.61192719
  -1.34149673]
 \begin{bmatrix} 0.47689837 & 0.14844958 & 0.52904524 & \dots & 0.0960042 & -0.0451133 \end{bmatrix}
   0.07912172]
 [ 0.85053068 -0.83912419 -1.01177408 ... -0.07223876  0.31112445
  -1.07836109]]
[[ 0.003193 ]
 [-0.01080922]
 [-0.01045412]]
[[ 1.06583858 -1.06149724 -1.79864091]]
[[0.15552509]]
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 13-neural_network.py

Check your code? QA Review

14. Train NeuralNetwork \(\text{\text{F}} \)

mandatory

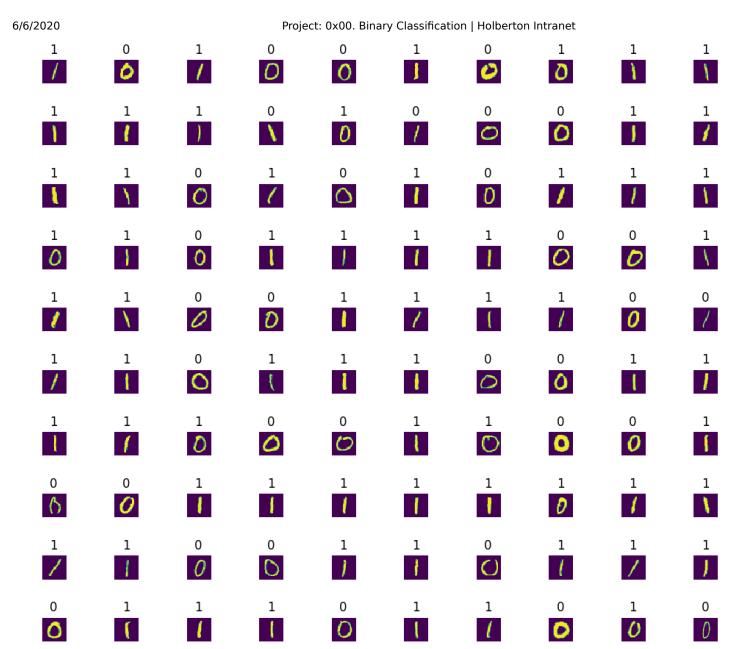
Score: 100.00% (Checks completed: 100.00%)

Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 13-neural_network.py):



- Add the public method def train(self, X, Y, iterations=5000, alpha=0.05):
 - Trains the neural network
 - X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
 - o alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
 - All exceptions should be raised in the order listed above
 - Updates the private attributes __w1 , __b1 , __A1 , __w2 , __b2 , and __A2
 - You are allowed to use one loop
 - o Returns the evaluation of the training data after iterations of training have occurred

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 14-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
NN = __import__('14-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
nn = NN(X_{train.shape}[0], 3)
A, cost = nn.train(X_train, Y_train, iterations=100)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = nn.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./14-main.py
Train cost: 0.4680930945144984
Train accuracy: 84.69009080142123%
Dev cost: 0.45985938789496067
Dev accuracy: 86.52482269503547%
alexa@ubuntu-xenial:0x00-binary_classification$
```



Pretty good... but there are still some incorrect labels. We need more data to see why...

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 14-neural_network.py

QA Review Check your code?

15. Upgrade Train NeuralNetwork mandatory

Score: 60.71% (Checks completed: 50.00%)

Write a class NeuralNetwork that defines a neural network with one hidden layer performing binary classification (based on 14-neural_network.py):

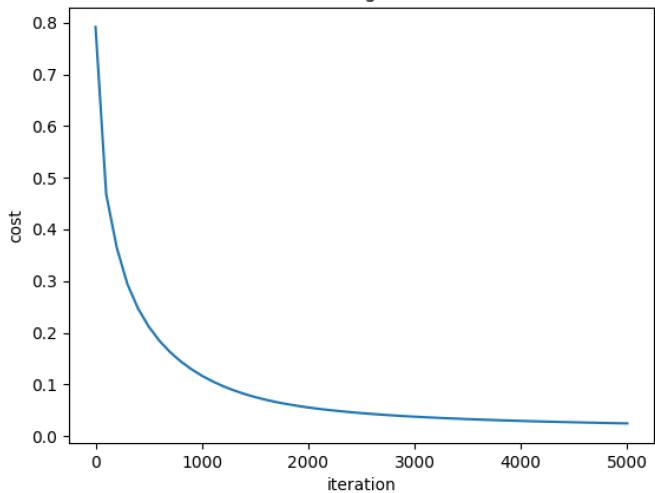
• Update the public method train to def train(self, X, Y, iterations=5000, alpha=0.05, verbose=True, graph=True, step=100):



- o Trains the neural network
- X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
- Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
- iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
- o alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
- Updates the private attributes __W1 , __b1 , __A1 , __W2 , __b2 , and __A2
- verbose is a boolean that defines whether or not to print information about the training. If True,
 print Cost after {iteration} iterations: {cost} every step iterations:
 - Include data from the 0th and last iteration
- graph is a boolean that defines whether or not to graph information about the training once the training has completed. If True:
 - Plot the training data every step iterations as a blue line
 - Labelthe x-axis as iteration
 - Label the y-axis as cost
 - Title the plot Training Cost
 - Include data from the 0th and last iteration
- o Only if either verbose or graph are True:
 - if step is not an integer, raise a TypeError with the exception step must be an integer
 - if step is not positive and less than or equal to iterations, raise a ValueError with the exception step must be positive and <= iterations</p>
- All exceptions should be raised in the order listed above
- The 0th iteration should represent the state of the neuron before any training has occurred
- You are allowed to use one loop
- Returns the evaluation of the training data after iterations of training have occurred

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 15-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
NN = __import__('15-neural_network').NeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
nn = NN(X_{train.shape}[0], 3)
A, cost = nn.train(X_train, Y_train)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = nn.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./15-main.py
Cost after 0 iterations: 0.7917984405648547
Cost after 100 iterations: 0.4680930945144984
Cost after 5000 iterations: 0.024369225667283875
```

Training Cost



Train cost: 0.024369225667283875 Train accuracy: 99.3999210422424% Dev cost: 0.020330639788072768 Dev accuracy: 99.57446808510639%



- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 15-neural_network.py



16. DeepNeuralNetwork mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification:

- class constructor: def __init__(self, nx, layers):
 - o nx is the number of input features
 - If nx is not an integer, raise a TypeError with the exception: nx must be an integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be a positive integer
 - o layers is a list representing the number of nodes in each layer of the network
 - If layers is not a list, raise a TypeError with the exception: layers must be a list of positive integers
 - The first value in layers represents the number of nodes in the first layer, ...
 - If the elements in layers are not all positive integers, raise a TypeError with the exception layers must be a list of positive integers
 - All exceptions should be raised in the order listed above
 - Sets the public instance attributes:
 - L: The number of layers in the neural network.
 - cache: A dictionary to hold all intermediary values of the network. Upon instantiation, it should be set to an empty dictionary.
 - weights: A dictionary to hold all weights and biased of the network. Upon instantiation:
 - The weights of the network should be initialized using the He et al. method and saved in the weights dictionary using the key W{1} where {1} is the hidden layer the weight belongs to
 - The biases of the network should be initialized to 0's and saved in the weights dictionary using the key b{1} where {1} is the hidden layer the bias belongs to
 - You are allowed to use one loop

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 16-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('16-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
print(deep.cache)
print(deep.weights)
print(deep.L)
deep.L = 10
print(deep.L)
alexa@ubuntu-xenial:0x00-binary_classification$ ./16-main.py
{}
{'b3': array([[0.]]), 'W2': array([[ 0.4609219 , 0.56004008, -1.2250799 , -0.09454199,
0.57799141],
       [-0.16310703, 0.06882082, -0.94578088, -0.30359994, 1.15661914],
       [-0.49841799, -0.9111359, 0.09453424, 0.49877298, 0.75503205]]), 'W3': array
([[-0.42271877, 0.18165055, 0.4444639 ]]), 'b2': array([[0.],
       [0.],
       [0.]]), 'W1': array([[ 0.0890981 , 0.02021099, 0.04943373, ..., 0.02632982,
         0.03090699, -0.06775582],
       [ 0.02408701, 0.00749784, 0.02672082, ..., 0.00484894,
        -0.00227857, 0.00399625],
       [ 0.04295829, -0.04238217, -0.05110231, ..., -0.00364861,
         0.01571416, -0.05446546],
       [ 0.05361891, -0.05984585, -0.09117898, ..., -0.03094292,
        -0.01925805, -0.06308145],
       [-0.01667953, -0.04216413, 0.06239623, ..., -0.02024521,
        -0.05159656, -0.02373981]]), 'b1': array([[0.],
       [0.],
       [0.],
       [0.],
       [0.]])
3
10
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 16-deep_neural_network.py

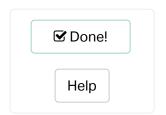
Check your code?

QA Review

17. Privatize DeepNeuralNetwork

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 16-deep_neural_network.py):

- class constructor: def __init__(self, nx, layers):
 - nx is the number of input features
 - If nx is not an integer, raise a TypeError with the exception: nx must be an integer
 - If nx is less than 1, raise a ValueError with the exception: nx must be a positive integer
 - layers is a list representing the number of nodes in each layer of the network
 - If layers is not a list, raise a TypeError with the exception: layers must be a list of positive integers
 - The first value in layers represents the number of nodes in the first layer, ...
 - If the elements in layers are not all positive integers, raise a TypeError with the exception layers must be a list of positive integers
 - o All exceptions should be raised in the order listed above
 - Sets the **private** instance attributes:
 - __L : The number of layers in the neural network.
 - __cache : A dictionary to hold all intermediary values of the network. Upon instantiation, it should be set to an empty dictionary.
 - weights: A dictionary to hold all weights and biased of the network. Upon instantiation:
 - The weights of the network should be initialized using the He et al. method and saved in the __weights dictionary using the key W{1} where {1} is the hidden layer the weight belongs to
 - The biases of the network should be initialized to 0 's and saved in the __weights dictionary using the key b{1} where {1} is the hidden layer the bias belongs to
 - Each private attribute should have a corresponding getter function (no setter function).
 - You are allowed to use one loop

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 17-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('17-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
print(deep.cache)
print(deep.weights)
print(deep.L)
deep.L = 10
print(deep.L)
alexa@ubuntu-xenial:0x00-binary_classification$ ./17-main.py
{}
{'b1': array([[0.],
       [0.],
       [0.],
       [0.],
       [0.]]), 'b2': array([[0.],
       [0.],
       [0.]]), 'W2': array([[ 0.4609219 , 0.56004008, -1.2250799 , -0.09454199, 0.577
99141],
       [-0.16310703, 0.06882082, -0.94578088, -0.30359994,
                                                             1.15661914],
                                                             0.75503205]]), 'W1': array
       [-0.49841799, -0.9111359 , 0.09453424, 0.49877298,
([[0.0890981, 0.02021099, 0.04943373, ..., 0.02632982,
         0.03090699, -0.06775582],
       [ 0.02408701, 0.00749784,
                                  0.02672082, ..., 0.00484894,
        -0.00227857, 0.00399625],
       [ 0.04295829, -0.04238217, -0.05110231, ..., -0.00364861,
         0.01571416, -0.05446546],
       [ 0.05361891, -0.05984585, -0.09117898, ..., -0.03094292,
        -0.01925805, -0.06308145],
       [-0.01667953, -0.04216413, 0.06239623, ..., -0.02024521,
        -0.05159656, -0.02373981]]), 'b3': array([[0.]]), 'W3': array([[-0.42271877,
0.18165055, 0.4444639 ]])}
3
Traceback (most recent call last):
  File "./17-main.py", line 16, in <module>
    deep.L = 10
AttributeError: can't set attribute
alexa@ubuntu-xenial:0x00-binary_classification$
```

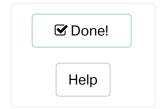
- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 17-deep_neural_network.py

Check your code?	QA Review

18. DeepNeuralNetwork Forward Propagation

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 17-deep_neural_network.py):

- Add the public method def forward_prop(self, X):
 - Calculates the forward propagation of the neural network
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Updates the private attribute __cache:
 - The activated outputs of each layer should be saved in the __cache dictionary using the key A{1} where {1} is the hidden layer the activated output belongs to
 - X should be saved to the cache dictionary using the key A0
 - All neurons should use a sigmoid activation function
 - You are allowed to use one loop
 - Returns the output of the neural network and the cache, respectively

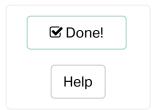
```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 18-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('18-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
deep._DeepNeuralNetwork__weights['b1'] = np.ones((5, 1))
deep._DeepNeuralNetwork__weights['b2'] = np.ones((3, 1))
deep._DeepNeuralNetwork__weights['b3'] = np.ones((1, 1))
A, cache = deep.forward_prop(X)
print(A)
print(cache)
print(cache is deep.cache)
print(A is cache['A3'])
alexa@ubuntu-xenial:0x00-binary_classification$ ./18-main.py
[[0.75603476 0.7516025 0.75526716 ... 0.75228888 0.75522853 0.75217069]]
{'A1': array([[0.4678435 , 0.64207147, 0.55271425, ..., 0.61718097, 0.56412986,
        0.72751504],
       [0.79441392, 0.87140579, 0.72851107, ..., 0.8898201 , 0.79466389,
        0.822570681,
       [0.72337339, 0.68239373, 0.63526533, ..., 0.7036234 , 0.7770501 ,
        0.694653461,
       [0.65305735, 0.69829955, 0.58646313, ..., 0.73949722, 0.52054315,
        0.73151973],
       [0.67408798, 0.69624537, 0.73084352, ..., 0.70663173, 0.76204175,
        0.72705428]]), 'A3': array([[0.75603476, 0.7516025 , 0.75526716, ..., 0.7522888
8, 0.75522853,
        0.75217069]]), 'A0': array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]
       [0., 0., 0., ..., 0., 0., 0.]
       [0., 0., 0., \ldots, 0., 0., 0.]
       [0., 0., 0., \ldots, 0., 0., 0.]
       [0., 0., 0., ..., 0., 0., 0.]], dtype=float32), 'A2': array([[0.75067742, 0.7831
9533, 0.77755571, ..., 0.77891002, 0.75847839,
        0.78517215],
       [0.70591081, 0.71159364, 0.7362214 , ..., 0.70845465, 0.72133875,
        0.71090691],
       [0.72032379, 0.69519095, 0.72414599, ..., 0.70067751, 0.71161433,
        0.70420437]])}
True
True
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 18-deep neural network.pv

Check your code? **QA Review**

19. DeepNeuralNetwork Cost mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 18-deep_neural_network.py):

- Add the public method def cost(self, Y, A):
 - Calculates the cost of the model using logistic regression
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - A is a numpy.ndarray with shape (1, m) containing the activated output of the neuron for each example
 - To avoid division by zero errors, please use 1.0000001 A instead of 1 A
 - Returns the cost

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 19-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('19-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
A, _{-} = deep.forward_prop(X)
cost = deep.cost(Y, A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./19-main.py
0.6958649419170609
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 19-deep_neural_network.py

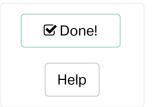
Check your code?

QA Review

20. Evaluate DeepNeuralNetwork

mandatory

Score: 100.00% (Checks completed: 100.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 19-deep_neural_network.py):

- Add the public method def evaluate(self, X, Y):
 - Evaluates the neural network's predictions
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - Returns the neuron's prediction and the cost of the network, respectively
 - The prediction should be a numpy.ndarray with shape (1, m) containing the predicted labels for each example
 - The label values should be 1 if the output of the network is >= 0.5 and 0 otherwise

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 20-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('20-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
A, cost = deep.evaluate(X, Y)
print(A)
print(cost)
alexa@ubuntu-xenial:0x00-binary_classification$ ./20-main.py
[[1 \ 1 \ 1 \ \dots \ 1 \ 1 \ 1]]
0.6958649419170609
alexa@ubuntu-xenial:0x00-binary_classification$
```

Repo:

• GitHub repository: holbertonschool-machine_learning



- Directory: supervised_learning/0x00-binary_classification
- File: 20-deep_neural_network.py

Check your code? **QA Review**

21. DeepNeuralNetwork Gradient Descent mandatory

Score: 0.00% (Checks completed: 0.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 20-deep_neural_network.py):

- Add the public method def gradient_descent(self, Y, cache, alpha=0.05):
 - Calculates one pass of gradient descent on the neural network
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - cache is a dictionary containing all the intermediary values of the network
 - alpha is the learning rate
 - Updates the private attribute __weights
 - You are allowed to use one loop

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 21-main.py
#!/usr/bin/env python3
import numpy as np
Deep = __import__('21-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_3D, Y = lib_train['X'], lib_train['Y']
X = X_3D.reshape((X_3D.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X.shape[0], [5, 3, 1])
A, cache = deep.forward_prop(X)
deep.gradient_descent(Y, cache, 0.5)
print(deep.weights)
alexa@ubuntu-xenial:0x00-binary_classification$ ./21-main.py
{'b3': array([[0.00659936]]), 'b2': array([[-0.00055419],
       [ 0.00032369],
       [ 0.0007201 ]]), 'W2': array([[ 0.4586347 , 0.55968571, -1.22435332, -0.0951687
   0.57668454],
4,
       [-0.16209305, 0.06902405, -0.9460547, -0.30329296, 1.15722071],
       [-0.49595566, -0.91068385, 0.09382566, 0.49948968, 0.75647764]]), 'b1': array
([[-1.01835520e-03],
       [-1.22929756e-04],
       [ 9.25521878e-05],
       [ 1.07730873e-04],
       [ 2.29014796e-04]]), 'W3': array([[-0.41262664, 0.18889024, 0.44717929]]), 'W
1': array([[ 0.0890981 , 0.02021099, 0.04943373, ..., 0.02632982,
         0.03090699, -0.06775582],
       [ 0.02408701, 0.00749784, 0.02672082, ..., 0.00484894,
        -0.00227857, 0.00399625],
       [ 0.04295829, -0.04238217, -0.05110231, ..., -0.00364861,
         0.01571416, -0.05446546],
       [ 0.05361891, -0.05984585, -0.09117898, ..., -0.03094292,
        -0.01925805, -0.06308145],
       [-0.01667953, -0.04216413, 0.06239623, \ldots, -0.02024521,
        -0.05159656, -0.02373981]])}
alexa@ubuntu-xenial:0x00-binary_classification$
```

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 21-deep_neural_network.py

Check your code?

Ask a new correction

QA Review

22. Train DeepNeuralNetwork

mandatory

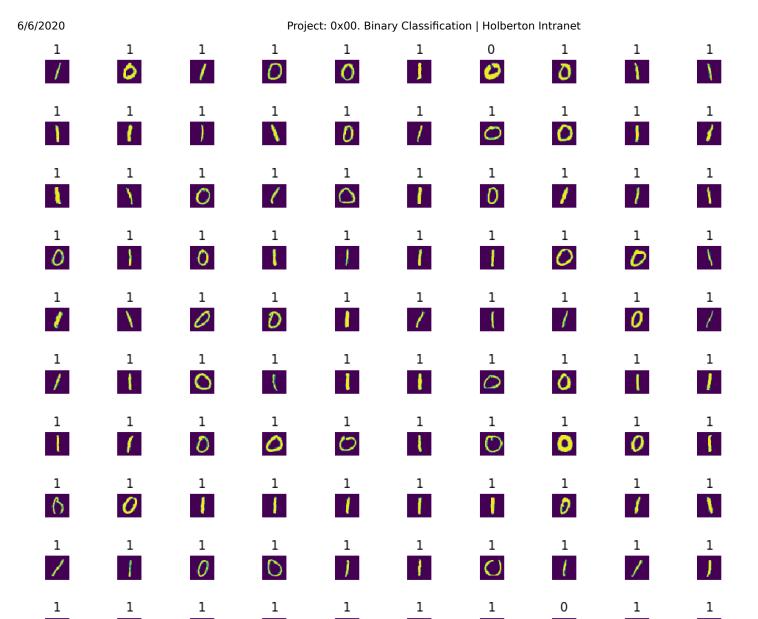
Score: 0.00% (Checks completed: 0.00%)



Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 21-deep_neural_network.py):

- Add the public method def train(self, X, Y, iterations=5000, alpha=0.05):
 - Trains the deep neural network
 - o X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - o Y is a numpy.ndarray with shape (1, m) that contains the correct labels for the input data
 - iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
 - o alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
 - All exceptions should be raised in the order listed above
 - Updates the private attributes __weights and __cache
 - You are allowed to use one loop
 - Returns the evaluation of the training data after iterations of training have occurred

```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 22-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
Deep = __import__('22-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X_train.shape[0], [5, 3, 1])
A, cost = deep.train(X_train, Y_train, iterations=100)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = deep.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./22-main.py
Train cost: 0.6444304786060048
Train accuracy: 56.241610738255034%
Dev cost: 0.6428913158565179
Dev accuracy: 57.730496453900706%
```



Hmm... doesn't seem like this worked very well. Could it be because of our architecture or that it wasn't trained properly? We need to see more information...

Repo:

- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification
- File: 22-deep_neural_network.py



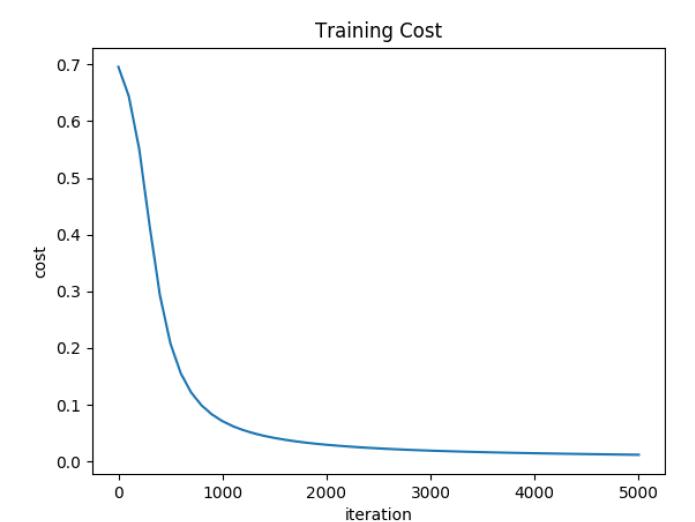
23. Upgrade Train DeepNeuralNetwork mandatory

Score: 0.00% (Checks completed: 0.00%)

Write a class DeepNeuralNetwork that defines a deep neural network performing binary classification (based on 22-deep_neural_network.py):

- ☐ Done? Help
- Update the public method train to def train(self, X, Y, iterations=5000, alpha=0.05, verbose=True, graph=True, step=100):
 - Trains the deep neural network by updating the private attributes __weights and __cache
 - X is a numpy.ndarray with shape (nx, m) that contains the input data
 - nx is the number of input features to the neuron
 - m is the number of examples
 - Y is a numpy ndarray with shape (1, m) that contains the correct labels for the input data
 - iterations is the number of iterations to train over
 - if iterations is not an integer, raise a TypeError with the exception iterations must be an integer
 - if iterations is not positive, raise a ValueError with the exception iterations must be a positive integer
 - alpha is the learning rate
 - if alpha is not a float, raise a TypeError with the exception alpha must be a float
 - if alpha is not positive, raise a ValueError with the exception alpha must be positive
 - verbose is a boolean that defines whether or not to print information about the training. If True, print Cost after {iteration} iterations: {cost} every step iterations:
 - Include data from the 0th and last iteration
 - graph is a boolean that defines whether or not to graph information about the training once the training has completed. If True:
 - Plot the training data every step iterations as a blue line
 - Label the x-axis as iteration
 - Label the y-axis as cost
 - Title the plot Training Cost
 - Include data from the 0th and last iteration
 - o Only if either verbose or graph are True:
 - if step is not an integer, raise a TypeError with the exception step must be an integer
 - if step is not positive and less than or equal to iterations, raise a ValueError with the exception step must be positive and <= iterations
 - All exceptions should be raised in the order listed above
 - The 0th iteration should represent the state of the neuron before any training has occurred
 - You are allowed to use one loop
 - Returns the evaluation of the training data after iterations of training have occurred

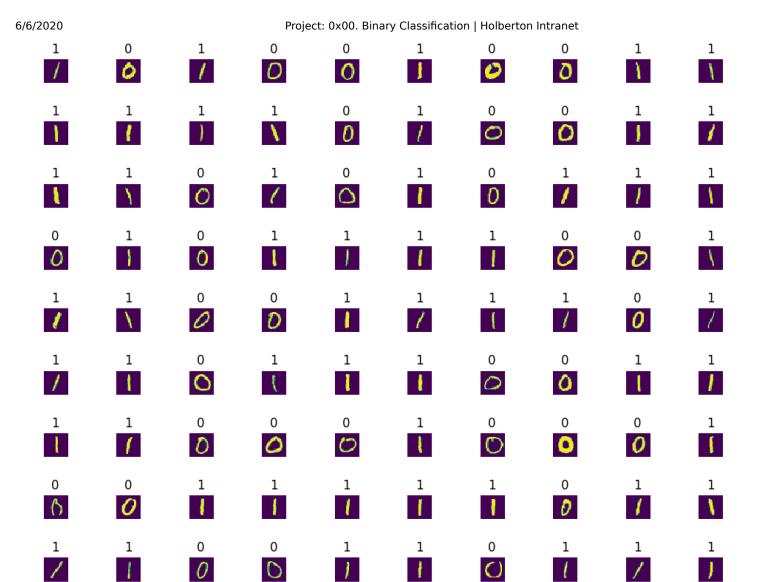
```
alexa@ubuntu-xenial:0x00-binary_classification$ cat 23-main.py
#!/usr/bin/env python3
import matplotlib.pyplot as plt
import numpy as np
Deep = __import__('23-deep_neural_network').DeepNeuralNetwork
lib_train = np.load('../data/Binary_Train.npz')
X_train_3D, Y_train = lib_train['X'], lib_train['Y']
X_{train} = X_{train_3D.reshape((X_{train_3D.shape[0], -1)).T}
lib_dev = np.load('../data/Binary_Dev.npz')
X_dev_3D, Y_dev = lib_dev['X'], lib_dev['Y']
X_{dev} = X_{dev}.shape((X_{dev}.shape[0], -1)).T
np.random.seed(0)
deep = Deep(X_train.shape[0], [5, 3, 1])
A, cost = deep.train(X_train, Y_train)
accuracy = np.sum(A == Y_train) / Y_train.shape[1] * 100
print("Train cost:", cost)
print("Train accuracy: {}%".format(accuracy))
A, cost = deep.evaluate(X_dev, Y_dev)
accuracy = np.sum(A == Y_dev) / Y_dev.shape[1] * 100
print("Dev cost:", cost)
print("Dev accuracy: {}%".format(accuracy))
fig = plt.figure(figsize=(10, 10))
for i in range(100):
    fig.add_subplot(10, 10, i + 1)
    plt.imshow(X_dev_3D[i])
    plt.title(A[0, i])
    plt.axis('off')
plt.tight_layout()
plt.show()
alexa@ubuntu-xenial:0x00-binary_classification$ ./23-main.py
Cost after 0 iterations: 0.6958649419170609
Cost after 100 iterations: 0.6444304786060048
Cost after 5000 iterations: 0.011671820326008168
```



Train cost: 0.011671820326008168 Train accuracy: 99.88945913936044% Dev cost: 0.00924955213227925

Dev accuracy: 99.95271867612293%

2



- GitHub repository: holbertonschool-machine_learning
- Directory: supervised_learning/0x00-binary_classification

1

• File: 23-deep_neural_network.py

1

1

QA Review Ask a new correction

O

1

Copyright © 2020 Holberton School. All rights reserved.

0



0

0